

Three training runs per configuration, seed the only difference

Solvability moves a lot

Validation loss barely moves

